

Contents

1	Getting Started	5
1.1	Setup	5
1.2	Designing a Nodal System	5
1.3	Sequence design with Web DD	6
1.4	Thermodynamic analysis and simulation with NUPACK	7
1.5	Next steps	7
2	Overview	9
2.1	Design Goals	9
2.2	Architecture	10
2.3	Contents of your Installation	10
2.4	Installing Workbench	11
2.5	Getting started	11
2.6	Further Reading	11
3	File Management	13
3.1	Creating files	13
3.2	Uploading files	13
3.3	Downloading files	14
3.4	Renaming files	14
3.5	Deleting files	14
3.6	Advanced information	14
3.6.1	Home directory	14
3.6.2	Packages	15

4	Workbench Applications	17
4.1	Behavioral Design	17
4.2	Sequence Design	17
4.3	Simulation and Analysis	18
4.4	Utilities	18
5	Behavioral Design	19
6	Designing Nodal Systems	21
6.1	Overview	21
6.2	Adding Nodes	22
6.3	Defining Motifs	22
6.3.1	Using the Motif Editor to edit or define motifs.	23
6.4	Complementarities	23
6.5	Errors	24
6.6	Compiling	24
7	Pepper Systems	27
7.1	Overview	27
7.2	Pepper File syntax	27
7.2.1	Pepper Components	27
7.2.2	Pepper Systems	31
7.3	Pepper Intermediate Language (PIL)	32
8	DyNAMiC Intermediate Language (DIL)	33
8.1	Viewing and Editing systems	33
8.1.1	Ribbon	34
8.1.2	Complexes pane	34
8.1.3	Segments pane	34
8.1.4	Strands pane	35
8.1.5	Interactivity	35
8.2	Designing sequences	35
8.3	Enumerating reactions	36
8.4	Evaluating thermodynamics	36

<i>CONTENTS</i>	3
9 Working with Sequences	37
9.1 Sequence Editor	37
9.2 Sequence Designers	37
9.3 Analysis and Simulation	38
10 Sequence Editor	39
11 Thermodynamic simulations with NUPACK	41
11.1 MFE Complexes	41
11.2 Pairwise Complexes	41
12 Simulation and Analysis	43
12.1 NUPACK partition function calculation	43
12.2 Mfold partition function calculation and MFE structure determination	43
12.3 Vienna RNAfold partition function calculation and MFE structure determination	43
13 Domain-level reaction enumeration	45
13.1 Preparing input for the enumerator	45
13.2 Running the enumerator	47
13.3 Viewing enumeration results	48
13.3.1 Enumeration results toolbar	49
13.3.2 Viewing complex and reaction details	50
13.3.3 Viewing full and condensed reactions	51
14 Thermodynamic simulations with NUPACK	55
14.1 MFE Complexes	55
14.2 Pairwise Complexes	55
15 Customization	57
15.1 Console	57
15.2 Custom scripts	57

16 DyNAMiC Workbench Server	59
16.1 Overview	59
16.2 Server Tools	59
16.3 Hosting your own server	60
16.3.1 Running the server	60
16.3.2 Logging in to the server	61
16.3.3 Using the web interface	61
16.3.4 Interacting with the server via SSH	62
16.3.5 Shutting down the server	62
16.3.6 Server Manager Application	63
17 Scripting	65
17.1 Overview	65
17.2 Console	65
17.3 Scripting Applications	65
18 Workbench Application Development	67
18.1 Overview	67
18.2 Bundling	67
18.3 API Documentation	67

Chapter 1

Getting Started

This document will get you started with Workbench as quickly as possible. For a more detailed introduction, see the [Overview](#).

1.1 Setup

Note: This part assumes you’ve already [installed](#) Workbench Server, or that you’re using a hosted installation. If you didn’t configure this server, you’re on a hosted installation; go on to the next section. See the [server](#) page if you’re confused.

1.2 Designing a Nodal System

To design a system using the Nodal Abstraction described in [Yin et al., 2008](#):

1. Create a new file
 - Select “New” from the main menu bar
 - Enter a filename, like “system.nodal” (you can omit the file extension)
 - From the menu under “Create,” select “Nodal System.” Your new file should appear in the tree on the left (labeled “Files”)
 - Open the file by double-clicking on its name in the Files tree
2. Add some nodes to the workspace
 - Drag and drop motifs from the palette labelled “Motifs” onto the main white area (the “Workspace”). Nodes can be renamed by clicking on the gray label above.

- For a full reference of the different motifs, see [nodal systems](#).
3. Connect nodes together
 - To indicate complementarities, select the Complementarity tool (labeled “Connect”). Click and drag from output ports (circular) to input ports (triangular).
 - You can delete complementarities by switching back to the pointer tool (arrow), and pressing the large red “X”
 4. Build your project and send output to a sequence designer
 - Select the “Build” tab, then click “Compile”. You should see a bunch of output files appear in the files tree to the left:
 - (nodal system name).txt: this is a serialized (textual) representation of the reaction graph you’ve just drawn
 - (nodal system name).svg: a graphical representation of the strands necessary to implement your system
 - (nodal system name).nupack: a script for the [NUPACK Multiobjective designer](#) to design sequences using [thermodynamic ensemble defect minimization](#).
 - (nodal system name).domains: a file which will allow you to quickly and easily generate and optimize sequences stochastically in your browser, using the Molecular Systems lab’s Web Domain Designer (Web DD).

[Read more](#) about the nodal designer.

1.3 Sequence design with Web DD

To design some sequences for the structure you’ve just described:

1. Open the .domains file by double-clicking its name in the Files tree
2. Click “Mutate” to begin optimizing the design. You’ll see the individual domains being designed in the center pane, a schematic of the various construct strands in the “Structure” pane to the right, and a real-time visualization of the candidate sequences in the “Strands” pane below. You can pause the mutation and save any part of this view by clicking the “Save” button and selecting a file name.
 - Mutate the design for a while
 - Once you’re satisfied with the score (lower is better), pause mutations by clicking the “Mutate” button again.
 - Save your final strands by clicking “Save” in the “Strands” pane, and selecting a file (how about “system.seq”).

[Read more](#) about Web DD, or [Sequence design](#).

1.4 Thermodynamic analysis and simulation with NUPACK

To perform full physical model calculations using NUPACK on the sequences we've just designed:

1. In the “Strands” pane within Web DD, select “Compute”
2. Hover over “Pairwise MFE Complexes,” then enter a file name and click “Run.” You don't need to provide an extension, just a prefix, since NUPACK will generate lots of files, and Workbench will wrap them up in a “package” file for you.
3. The Console should pop up, and a message should be displayed indicating that the NUPACK task is being run. Wait until you see a bunch more output (starting with “Permutation generation complete.”). This will indicate that the task has completed. Your Files tree will also refresh.
4. You should see a new file called (file name you entered earlier).package; double-click it to view the results, which may take a moment to load. You should see a list of the distinct strands analyzed, a graph of the minimum free energy (G) for the various possible complexes, and a list of complexes ordered by concentration. Click one of these complexes in order to view the minimum free energy structure as a 2D visualization, an arc diagram, or a matrix visualization.

[Read more](#) about using NUPACK in Workbench, or check out the [NUPACK Website at Caltech](#)

1.5 Next steps

- Read more about the [Applications](#) available in Workbench
- Learn more about the [File system](#)
- Discover how to [customize Workbench](#) by scripting and developing applications.

Chapter 2

Overview

DyNAMiC Workbench is an *Integrated Development Environment* for dynamic DNA systems. This document aims to survey the goals and architecture of Workbench and to acquaint you with important features. For a more rapid introduction, see the [Getting Started guide](#).

2.1 Design Goals

Workbench has been designed with several goals in mind:

- **Integration** - Workbench allows access to the full set of tools necessary for designing assembly and computation systems based on DNA strand displacement, and it streamlines the workflow between these tools.
- **Usability** - Workbench presents an intuitive, cross-browser/cross-platform, web-based interface
- **Scalability** - Workbench has built-in task deployment system (see *Architecture* below) which can execute computationally-intensive tasks on a variety of targets, including the local web server, a cluster such as Orchestra, or a number of remote web services, such as the NUPACK or Mfold web servers.
- **Extensibility** - Workbench includes powerful, but simple tools for automating various tasks. The client-side interface is entirely scriptable, and tools are included for developing entirely new client-side applications. New server tools can be easily added to Workbench as well. Since the client and server are written in the same language (Javascript), code can be easily re-used between the client and server.

2.2 Architecture

Workbench includes two main components:

- **Workbench Server** – The server is responsible for storing and managing files, as well as deploying and executing computationally intensive tasks. If you’re accessing a hosted version of Workbench, you likely won’t interact directly with the Workbench server at all. If you’re hosting Workbench yourself, it’s helpful to understand a bit about the server. [Read more](#).
 - **Server tools** - The server can run several types of computational tasks, each of which is encapsulated by a *server tool*. Server tools require a wrapper written in Javascript, but they can easily access and call tools written in any language. Most tools included with Workbench are written in C, while others are written in Python. [Read more](#)
 - **File management** - The server maintains a “home directory” for each user. Many types of files can be opened and edited within Workbench, including DNA-related files (.seq, .nupack, .nodal), and other relevant files not strictly related to DNA (.txt, .xml, .html, .tex), while other file types (.svg, .pdf) can be previewed from within Workbench. [Read More](#)
 - **Users** - To access the server, you need a user account. This is true even if you’re hosting your own server instance. Workbench will redirect you to a login page from which you can make an account. [Read More](#)
- **Workbench Client** - The client is the main interface to Workbench. It can be run as a standalone application on Mac OS X, accessed within any modern web browser.
 - **Applications** - Workbench does lots of different things; *Workbench Applications* implement all of these different tasks on the client. Workbench includes applications for everything from [behavioral design with the nodal abstraction](#) to [stochastic sequence design with Web DD](#). You can even write your own Workbench applications [Read More](#)
 - **Tasks** - Workbench allows you to launch and manage server tools from the client using the Task Manager. [Read More](#)

2.3 Contents of your Installation

If you downloaded Workbench as a disk image (.dmg) file for Mac OS X, several utilities have been provided to make your life a bit easier. These applications will be copied to your Applications directory as part of the [installation](#) process

- DyNAMiC Workbench.app - This contains a “site-specific browser” – a mini web browser, automatically configured to connect to a locally hosted Workbench server. This is just for convenience; you can just as easily navigate to the specified URL in a web browser.
- Workbench Server.app - This is a helper to allow you to launch and log in to the server. To start the server:
 1. Launch the ‘Workbench Server’ application from your Applications folder
 2. Click ‘Start’ from the server control window. You will see a Terminal window open with lots of output. You’ll also see a VirtualBox window open. Wait until you see `webserver-user@192.168.56.10's password::`; then enter the password for your server (the default is " ", a single space).
 3. Wait until you see `Server running from /media/sf_vmshare/infomachine2 at http://192.168.56.10:3000`. Congratulations, your server is running!

See the [Getting Started guide](#) for details on how to start on another operating system.

2.4 Installing Workbench

See [Installation](#) for details on how to install Workbench server.

2.5 Getting started

See the [Getting Started guide](#) for a quick introduction to the various things you can do with Workbench.

2.6 Further Reading

- [Getting Started](#)
- [Applications](#)

Chapter 3

File Management

Workbench has a flexible system for storing and managing files, in order to make the round-trip workflow between various applications a lot easier.

3.1 Creating files

You can easily create files within the Workbench interface:

1. Select “New” from the main menu bar
2. Enter a filename (you can choose to omit the file extension; workbench will add it)
3. From the menu under “Create,” select the type of file you would like to create. Your new file should appear in the tree on the left (labeled “Files”). There are some file types which are not currently available in Workbench, but for which support is planned; these appear as greyed out menu items.
4. Open the file by double-clicking on its name in the Files tree, or right-clicking and selecting “Open” from the context menu.

3.2 Uploading files

To upload a file to Workbench from your desktop, just drag and drop the file from your normal file manager (Finder, Nautilus, Explorer, etc.) onto the Files tree. Your new file will appear once it’s been uploaded (you should see a notification).

3.3 Downloading files

To download a file, right click on the file name in the Files tree, and select “Download”; it will download as an attachment in a new window.

3.4 Renaming files

To rename a file, right click on the file name in the Files tree, and select the “Rename” field, where you can enter a new file name. As soon as you leave the field, your file will be renamed.

You can also drag and drop files in the tree in order to arrange them to different folders.

Note: Workbench creates some special files, which cannot be renamed. These files are used to store system information, such as preferences or the contents of `.package` files (see “Packages” below). You can edit these files if you wish, but you cannot rename them.

3.5 Deleting files

To delete a file, right click on the file name and select “Delete”. Files deleted by Workbench cannot be recovered, unless they are independently backed up

3.6 Advanced information

3.6.1 Home directory

Each [user](#) gets a home directory on the Workbench server. This directory corresponds to that user’s email address. On hosted versions of Workbench, that directory will generally be made available to users via SFTP or SSH. If you’re hosting your own version of the Workbench Server, the “files” directory (where all users’ files are stored) is available as a VirtualBox “[shared folder](#)” between the virtual machine where Workbench Server runs, and your host machine. During [Installation](#) you should be prompted for this location. Once you’ve used workbench a bit, you can access your files using your normal file manager (Finder, Nautilus, Windows Explorer), as well as managing them through the workbench user interface. Just navigate to your home directory within ([shared folder](#))/files/(your email address).

3.6.2 Packages

Packages are special folders which contain lots of files, but largely one useful set of data. They were introduced because some [server tools](#), e.g. [NUPACK](#), generate lots of files, which are sort of cryptic on their own, but much more useful when aggregated together into a single file. However, it may sometimes be useful to look or use the original file generated by the NUPACK executable; to address this, Workbench provides the following bargain:

- For certain server tools, the multi-file output is directed to a single folder.
- Contents of that folder are intelligently aggregated into a single file.
- The folder is given a “.package” extension (but remains a folder, so you can open and browse it like any other)
- A “Package contents” file is written to the folder; this allows you to double-click the package itself (from the Files tree) to see the aggregated version, or to expand the package like any other folder, and view the individual output files generated by the server tool.

There is another kind of package file, called an Application Bundle, which can be used for creating custom Workbench applications. See [customization](#) for details.

Chapter 4

Workbench Applications

4.1 Behavioral Design

- [Nodal Designer](#) - Use a graphical language to specify relations between behavioral nodes; automatically convert behaviors into implementation systems.
 - [Strand editor \(DIL\)](#) - Graphical interface for viewing and editing a segment-level description of a nucleic acid system. Use to view the results of a nodal build, and to generate sequences with a sequence designer.
- [Pepper Designer](#)
 - [Pepper intermediate language \(PIL\)](#)
- CRN Designer (coming soon)

4.2 Sequence Design

- [Web DD](#) - A stochastic, domain-based sequence designer; can be used to design sequences for large systems very quickly by designing a set of noninteracting domains and then threading those domains together to form full strands.
- [NUPACK](#) Thermodynamic sequence designer - Uses Caltech's NUPACK web server to perform multi-objective thermodynamic sequence design. Enter a design using the NUPACK multi-objective sequence design script, select relevant parameters, and click "Design"; the task will be submitted to the Caltech server, and a popup window will be opened taking you to the results page.
- [Multisubjective](#) sequence designer

4.3 Simulation and Analysis

See [Simulation and Analysis](#)

- [Reaction Enumerator](#) – Enumerates the possible reactions and intermediate complexes between a set of initial complexes, at domain resolution
- [NUPACK](#) partition function calculation - Uses Caltech’s NUPACK web server to compute the minimum free energy secondary structure of a strand or set of strands
- [Mfold](#) partition function calculation and MFE structure determination - Uses the University of Albany’s DINAMelt web server to compute the minimum free energy secondary structure and base pair probabilities.
- [TBI Vienna RNAfold](#) partition function calculation and MFE structure determination - Uses TBI Vienna’s RNAfold Websuite to compute the minimum free energy secondary structure of single DNA or RNA strands.

4.4 Utilities

- [Sequence editor](#)
- Segment Threader - Allows an arbitrary set of sequences to be threaded together into a strand. A set of named sequences can be entered using the pane on the left, and a set of strand specifications (e.g. `strand1 = a b c d b* a*`) can be entered on the right. Clicking the “Thread” button will produce a set of sequences on the bottom generated by applying the strand specification rules using the named sequences.
- Structure Editor - Accessed by “Tools” > Structure Editor from the main toolbar. Allows simple preview of a secondary structure entered in dot-paren notation.

Chapter 5

Behavioral Design

Workbench comes with two behavioral designers:

- [Nodal](#) - The nodal designer allows users to graphically design systems using the nodal formalism described by Yin et al. 2008, and focuses on behavior and geometry.
- [Pepper](#) - The Pepper compiler uses a powerful, text-based input formal for designing arbitrary strand-displacement systems

Chapter 6

Designing Nodal Systems

6.1 Overview

The nodal formalism allows you to express complicated computational or assembly processes in terms of simple behavioral units, called nodes. Many types of nodes exist—these types are called “motifs”; motifs are defined by mapping a structural unit of DNA or RNA (such as a hairpin) to a simple behavioral function. Nodes generally have several *ports* which correspond to *domains* of the underlying nucleic acid species. Nodal programs are written by adding instances of motifs (nodes) to a workspace, then connecting the ports together to indicate behavioral relationships. These behavioral relationships in turn imply sequence complementarities. The nodal compiler, called DyNAMiC (the Dynamic Nucleic Acid Mechanism Compiler) propagates those sequence complementarity requirements to generate a list of distinct sequences which must be designed by a sequence designer (such as [DD](#) or [NUPACK](#)).

The basic workflow is like this:

- Nodal systems are assembled using the nodal designer
- The nodal compiler runs in real-time, verifying the system as you design it.
- Once your system has been designed you can generate a DIL (DyNAMiC Intermediate Language) file, which represents your system as a scheme of complementarity relationships between segments of the strands in your ensemble (see “Compiling”, below)
- From the DIL file, DyNAMiC can generate input files for many different sequence designers.
- Sequences can be designed using sequence designers

DyNAMiC actually uses a [JSON](#)-based input format, called [DyNAML](#). You

can enter DyNAML directly using the DyNAML editor application, or you can design custom DyNAML motifs from within workbench. The DyNAML language is described in [this whitepaper](#).

6.2 Adding Nodes

You can add nodes to the system by dragging and dropping existing motifs from the “Standard” panel in the lower-left, onto the workspace. You can also define new motifs (see [below](#)), which will appear in the “Custom” tab of the same panel.

6.3 Defining Motifs

You can create new motifs in two ways, each using the “Create Motif” tool in the ribbon. First, click the “Create Motif” button, then either:

- Click on the workspace to generate an empty motif. Then select this motif and use the inspector on the right to describe its implementation. You can either:
 - Click the “Edit Motif” button, and use a graphical interface to define the structure of the motif, or
 - Expand the “DyNAML Code” box and enter a custom [DyNAML](#) description. If you want to use one of the built-in motifs as template, click the “Copy from Built-in” button and select a built-in motif from the dropdown menu.
- Alternatively, click and drag on the workspace to select an existing system of nodes, and wrap that system in a new motif.
 - This will effectively remove those nodes from the system, and transform them into part of the motif.
 - All outgoing connections will be removed.
 - To “expose” ports within the motif to external complementarities, use the “Add Port” or “Expose” tool:
 - * “Add Port” : Select the “Add Port” tool, then click on the motif to add a port (hold ‘alt’ for an input, or ‘shift’ for a bridge). Then use the “Expose” tool to drag from the internal port to the external (motif-level) port
 - * “Expose”: Select the “Expose” tool; drag from the internal port to the motif itself; a new port will be created and exposed
 - You’ll need to instantiate the new motif by dragging and dropping it from the “Custom” panel in the lower-left.

If you decide you want to restore the nodes inside a motif to the workspace as normal nodes, select a motif and click the “Unwrap motif” button.

6.3.1 Using the Motif Editor to edit or define motifs.

The Motif Editor is a small graphical tool for defining custom motifs. It can be launched by creating a new motif (see above), selecting the motif in the workspace, and clicking “Edit Motif” in the inspector on the right side. To create a new motif:

- Start in the “Segments” pane in the top-left. Click the “Add” button to add new segments. You can click the arrow to the right to add segments of different lengths, or to add many specific sequences.
- Then use the “Strands” pane to thread segments together into strands. Click “Add” to add a new strand, and then type a strand specification in the DyNAML compact format. You can add as many strands as you’d like.
- In the “Structure” pane, click the “Strand Order” field and enter the names of your strands, separated by + signs, to describe the order in which the strands should appear in the starting complex. Once your cursor leaves the Strands field, the field will be populated with a simple graphic depicting the segments that comprise the strands you’ve selected. Note: If you create a strand using the “Strands” pane but don’t add the strand name to the “Strand Order” field, it will be omitted from the final motif.
- Then, in the “Structure” field, enter the structure for your strand in dot-parenthesis notation. Once a valid structure has been entered, an image of your complex and a depiction of the nodal motif will appear to the right.

6.4 Complementarities

To declare complementarities between ports:

- Select the “Complementarity” tool
- Drag from one port to another to declare complementarity. You may drag from Input to output, or Bridge to bridge; other connections will be disallowed. You may also add connections between ports within a motif, but not between ports inside and outside the motif.

Complementarities may only be drawn from output ports to input ports, or between bridge (square) ports. Some connections are not possible, depending on the shape, or “footprint” of the underlying domain. For instance, a connection between a domain with 3 segments: **a b c**, each of length 8, and a domain with

4 segments **a b c d** each of length 4 would be invalid, since the total number of nucleotides (24 and 16) do not match up. Likewise, a connection between a domain with 2 segments **a b** each of length 8 (total length 16) and a domain of one segment **a** of length 16 (total length 16) would *also* be invalid, since the number of segments in each domain is different. The nodal editor doesn't prohibit these connections from being drawn, but it will highlight the relevant nodes in red and report an error. See “Errors” below for how DyNAMiC and the Nodal editor handle these errors.

You can inspect the footprint of a port by moving your mouse over the port. A small tooltip will appear and will display the length of each segment in the underlying domain

6.5 Errors

Some connections may be drawn which are invalid. In these cases, DyNAMiC will report an error, and the Nodal interface will indicate this by changing the Build status indicator (in the lower right) to read “Error”. You can mouse over this field to view the error message. You can also click the field and select “Check for errors” to recompile the system, or “Show full results” to view the compiled library before sending to a sequence designer.

6.6 Compiling

Once your system is assembled, you can generate a picture of the underlying species, and generate input to sequence designers, using the “Build” tool. Select the “Build” tab in the ribbon, then click “Compile”. This will generate a new file with the extension `.dil`; this is a DyNAMiC Intermediate Language (DIL) file, and it contains your compiled system. The DIL file will open automatically, and you should see the species used to generate your system. From there, you can send results to one of several sequence designers. See [DIL system editor](#) for details.

Additionally, you can bypass the DIL step and select “Build all targets” from the dropdown next to the “Compile” button. This generates several files, which will appear as siblings to your system in the files tree:

- `(system).dynaml` - A formal, textual representation of your system (in the Dynamic Nucleic Acid Markup Language–DyNAML)
- `(system).svg` - A graphical representation of the underlying species in your system
- `(system).nupack` - A textual input file for the NUPACK multi-objective thermodynamic sequence designer

- (system).domains - A textual input file for [WebDD](#)
- (system).pil - A version of the compiled system in the [Pepper Intermediate Language](#) (PIL)

You can select several other options from the drop-down next to the “Compile” button:

- Clean – Deletes any of the above files in the same directory as your system.
- Compile with Compiler v2b – Compile with an old version of the compiler
- Compile locally – Compile in your browser (not on the server); this will not generate any output, but can be used for debugging.

Chapter 7

Pepper Systems

7.1 Overview

Pepper is a DNA circuit compiler developed by Shawn Ligocki in Erik Winfree's group at Caltech. It assists DNA programmers in building DNA computers by providing a templating language for specifying generic components and interfacing with state of the art designers and kinetic simulators to create and test sequences. Much like [DyNAMiC](#), it can be used to design generic components and to

Pepper systems are comprised of *components* and *systems*. Components are re-usable “building blocks,” which can be composed together to form systems. The rest of this documentation was written by Shawn Ligocki, and describes the syntax for Pepper component and system files.

7.2 Pepper File syntax

7.2.1 Pepper Components

A component file (name.comp) contains the specification for a single component in a DNA system (like an AND gate or a threshold, etc.). It specifies the secondary structure, kinetic and sequence constraints for that component. System files are used to connect together components.

The component syntax is based on Joe Zadeh's design specification syntax.

Comments

The pound sign (#) denotes that the rest of the line is a comment (python/sh style comments):

```
# This is a comment
```

Comments may appear on their own lines or after a command, like so:

```
declare component And22: x + y -> s + c # This is the function declaration
```

Component Declaration

In the spirit of Matlab's first line declarations, each component needs to have a declaration statement at the first line of code. Syntax:

```
declare component : ->
```

For example:

```
declare component And22: x + y -> s + c
```

The and are '+' separated lists of sequences that will be defined and used below in the specification. These sequences will be constrained on the system level and so usually represent the recognition regions of the inputs and outputs. Waste byproduct will generally not be represented here, because it will not be constrained on the system level.

should be the same as the base name of the component file (i.e. this file should be HalfAdder.comp).

Sequences

In order to constrain some regions/subdomains of some strands to be complementary to regions on other strands we define sequence regions:

```
sequence = :
```

For example:

```
sequence x = "6N S 13N S" : 21
```

The total length constraint is optional; if you don't want it, omit it and the colon.

We include the (redundant) length input for quick error checking as complex constraints can be easily misnumbered. In addition one wildcard (?) may be used instead of a number, to imply (make this as large as possible to fit the specified total length. For example:

```
sequence dx = "S ?N S" : 15
```

is the same as:

```
sequence dx = "S 13N S" : 15
```

Sequence region definitions may include previously defined regions, for example:

```
# The toehold, clamp and data region of the carry-bit output
sequence tc = "6N" : 6
sequence cc = "1S" : 1
sequence dc = "29N" : 29
# We combine them into a single label
sequence c = tc1 cc1 dc1 : 36
```

So now, for example:

```
c* = dc* cc* tc*
```

We use quotes for nucleotide constraints so that we can distinguish them from other sequence names and group them (each pair of quotes denotes one domain).

Strands

Strands represent individual strand of DNA in the system. The syntax is:

```
strand = :
```

For example:

```
strand C = "?N" c : 44
```

The “?N” is an explicit constraint. First of all, it uses the wildcard, so it fills the remaining space on the left. Furthermore, since it is never declared as a sequence it will be constrained only by what it is forced to bind to. The idea is that you have a data region of the strand and some other regions that you’re required to have, but don’t want to name.

Structures

Now we glue the strands together to make multi-stranded (or single-stranded) structures:

```
structure = :
```

So continuing the above example:

```
structure Gate = X + C + S + Y : U6 H15(+ H15(U29 + U14 H15(+)))
```

using [Joe Zadeh’s (H)elix (U)npaired notation](Joe Zadeh’s (H)elix (U)npaired notation “wikilink”) or, say

```
structure Gate = X + C + S + Y : 6. 15( + 15( 29. + 14. 15( + 45)
```

in a sort of shorthand dot-paren notation (where each number is a multiplier on the symbol following it).

By default it is assumed that you want to optimize the thermodynamics of the structure to be as close to the specified secondary structure as possible. If not, you can tell the compiler not to optimize a specific structure. For example:

```
structure [no-opt] Gate = X + C + S + Y : U6 H15(+ H15(U29 + U14 H15(+)))
```

This can be useful for two reasons. First, you may wish to impose base-pairing constraints but don't want thermodynamic optimization, for example, because you know this structure is not the MFE structure. Second, you may want the compiler to know that these strands should be grouped together as a complex, but the intended complex has pseudoknots that cannot (yet) be properly expressed in Pepper.

Alternatively for limited optimization:

```
structure [10nt] Gate = X + C + S + Y : U6 H15(+ H15(U29 + U14 H15(+)))
```

This aims to be within 10 nucleotides of the specified structure, on average. Note that how these parameters are interpreted is up to the back-end sequence designer. E.g. Zadeh's designer will try to minimize the average number of incorrect nucleotides; another designer might want the MFE structure to be within 10 nucleotides; the SpuriousC designer could use the parameter for some kind of weighting, but currently doesn't.

We can also specify the secondary structure on the domain level:

```
structure Gate = X + C + S + Y : domain .(+.(+. (+)))
```

assuming, for example, that the strands had been defined as

```
strand X = toe_x data_x
strand C = carry "29N"
strand S = "14N" sum
strand Y = sum* carry* data_x*
```

Here each dot-paren represents an entire domain.

Kinetics

Now all we have left is to explain the desired kinetics, what structures will interact and what will they produce:

```
kinetic ->
```

So if we're working with the half adder, we might have:

```
kinetic inX + Gate      -> waste_X + inter_G
kinetic inY + inter_Gate -> waste_Y + outS + outC
```

I might allow some optional parameters to fine tune these, maybe specifying fuzzy states or desired speed of reactions.

Examples

- [DNA compiler/And22.comp](#)

7.2.2 Pepper Systems

A system file (name.sys) contains a specification for the connectivity of [components](#) in a DNA system. It specifies each component and and ties their signal sequences together. A system may be used as a component in a larger system.

Comments

The pound sign (#) denotes that the rest of the line is a comment (python/sh style comments): # This is a comment Comments may appear on their own lines or after a command, like so:

```
import And22 # Half adder is used in the first layer only
```

System Declaration

In the spirit of Matlab's first line declarations, each system needs to have a declaration statement at the first line of code. Syntax:

```
declare system : ->
```

For example:

```
declare system HalfAdder: x0 + y0 -> s0 + c1
```

The and are '+' separated lists of sequences that will be constrained to components below.

These are used if this is a system being used as a component. However, for top-level systems, we will not have inputs and outputs and so

should be the same as the base name of the component file (i.e. this file should be HalfAdder.comp).

Imports

In order to use a component file you must import it first (python-style). Syntax:

```
import For example:
import And22
```

You may import from a different directory/name than the name you use in the file. For example:

```
import Georg0711/Parallel_And22 as And22
```

which would import from the component file ‘Parallel_And22.comp’ from the directory ‘Georg0711/’, but still use the name And22 in the specification.

Components

This is the meat of the circuit file. You must make one statement for each component in the DNA system specifying the input and output sequences. The syntax is:

```
component = : ->
```

for example

```
component G0_01 = And22: nx0 + y0 -> s0 + nc1
```

Note: Components **may** have 0 inputs or 0 outputs (or even both, if you can find that useful). Example:

```
component DNS0 = Detector: ns0 ->
```

Is a ‘ns0’ detector. It will activate if ns0 is input but doesn’t produce any outputs for downstream gates.

Examples

- [DNA compiler/Half Adder](#)
- [DNA compiler/Two-Bit Adder](#)

7.3 Pepper Intermediate Language (PIL)

Chapter 8

DyNAMiC Intermediate Language (DIL)

The DyNAMiC Intermediate Language (DIL) is the intermediate representation generated by the [Nodal compiler](#), DyNAMiC. DyNAMiC converts behavioral representations into molecular implementations using a set of pre-defined translation rules; these rules map “node types” to “molecule types.” The result of nodal compilation is therefore a segment-level representation of a set of complexes with a number of explicit Watson-Crick complementarity and orthogonality relationships, but no real nucleotide sequences.

The DIL Editor lets you visualize the primary and secondary structure of these complexes, specify sequence constraints for the segments which compose them, and export these structures to various sequence designers.

8.1 Viewing and Editing systems

To design a system from scratch, follow this procedure, guided by the figure:

1. In the *Segments* pane, click “Add” to add several segments (you can edit their names, sequences, or colors if you’d like)
2. In the *Strands* pane, click “Add” to add several strands.
3. Click the blank space under the “Segments” column (still in the Strands pane), and type a description of the strand. You can either use [DyNAML short notation](#) to group segments into domains, or you can just list the segments you wish to use, separated by spaces (e.g. 1 2 3* a 4 x, etc.); the “Sequence” column should be filled in automatically with the sequence you’ve just described

4. In the main *Complexes* pane, click “Add” to add a complex; select the new complex and click “Edit” or double-click the complex name to view a pop-up editor for the complex
5. Under “Strand Order”, click the box (which says “Click to add”), and type the names of each of the strands you want in the complex, *separated by + signs*.
6. Under “Structure”, enter a segment-wise structure for the complex, in dot-parenthesis notation. Once you have entered a valid structure, you should see a preview in the right-hand pane. You can close the complex editor popup once you’re done.

The following sections describe the interface in more detail.

8.1.1 Ribbon




Figure 8.1: Ribbon toolbar

Use the ribbon toolbar to **design sequences**, explore system **kinetics**, make **thermodynamic calculations**, or export the system to a Scalable Vector Graphics (SVG) file.

8.1.2 Complexes pane

The main interface pane displays a grid of complexes. Complexes are displayed as a secondary structure “planar graph.” A “View” menu in the lower-right corner of this panel may be used to modify the visualization; bubbles representing the individual nucleotides may be colored according to segment identity, domain identity, or base identity. Alternatively, the bubbles may be hidden and text labels may be used instead.

8.1.3 Segments pane

Shows the sequences of the segments which compose the system. Segments are discrete regions of sequence complementarity. A segment’s sequence can be edited by double clicking on the sequence and typing a new one; the complexes and strands panes will be updated automatically.

8.1.4 Strands pane

Shows the strands which comprise the system. Strands may not currently be edited directly, but their composite sequence and component segments/domains may be viewed.

8.1.5 Interactivity

You can mouse over segments in the Segments or Strands pane in order to highlight them elsewhere in the system. This can be useful for tracing where a particular segment has propagated within the system. The segment will be highlighted in yellow, and its complement will be highlighted in blue.

The screenshot shows the software interface with the following components:

- Top Bar:** Design Sequences (NUPACK, MS Multisubjective), Kinetics (Enumerate, to PIL, to SBML), Thermodynamics (Predict, MFE Structures, Pairwise), Export (SVG, DyNAML), and Workspace (Save, Help).
- Segments Pane (Right):**

Name	Sequence	Len
1	ATC>NNNN	8
2	NNN GGG NNN	9
3	NN	2
4	NNNNNNNN	8
5	NNNNNNNN	8
6	NNNNNNNN	8
7	NN	2
- Strands Pane (Bottom):**

Name	Sequence	Segments	Complex
n0_m0	...GAT...NNNNCCNNNN NN NNNNNNNN	A 1* 2* 3* 4*	
n1_m1	NNNN Segment: 1* / Base: 3 C NNNNNN NNN...	A 4 3 2 1 B 5* 6* 7* 1* C 2* 3*	
n2_m1	ATC NNNNNN NN NNNNNNNN NNNNNNNN NNNNNN...	A 1 7 6 5 B 4* 8* 9* 5* C 6* 7*	
n3_m1	NNNNNNNN NN NNNNNNNN NNNNNNNN NNNNNNNN...	A 5 9 8 4 B 1* 2* 3* 4* C 8* 9*	

Figure 8.2: Interactively highlight segments

8.2 Designing sequences

From the ribbon, DIL systems may be exported to various sequence designers, including [Web DD](#), [NUPACK](#), and [Multisubjective](#). Clicking these buttons will generate a file and open the requisite sequence designer.

8.3 Enumerating reactions

From the ribbon, DIL systems may be exported to a file readable by the [reaction enumerator](#). Simply click the “Enumerate” button, and an [enumerator](#) tab will open. Then click “Run enumerator” to run the reaction enumerator:

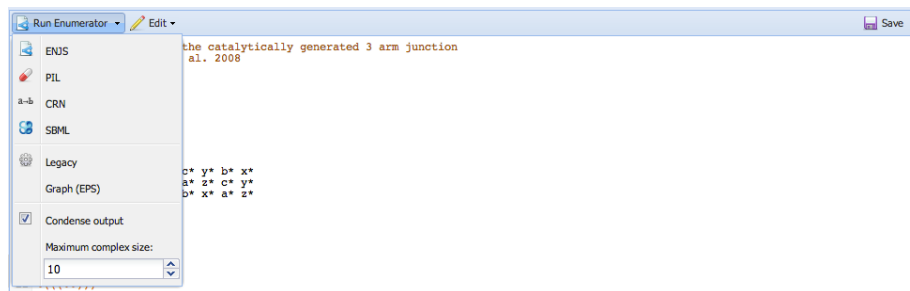


Figure 8.3: Run enumerator button and menu

See the [Enumerator](#) page for details.

8.4 Evaluating thermodynamics

From the ribbon, you can invoke various thermodynamic calculations. Click “Predict” to open a window for analysis using [NUPACK](#), or click the arrow to analyze using [Mfold](#) or [RNAfold](#). You can use the “MFE Structures” button to open a NUPACK analysis that will examine only single-stranded complexes (the “complex size” parameter is set to 1); the “Pairwise” button will open a NUPACK analysis with the “complex size” parameter set to 2.

Chapter 9

Working with Sequences

Workbench has many powerful tools for working with sequences of bases.

9.1 Sequence Editor

The sequence editor provides lots of tools for working with sequences as standard text files: The editor color-codes bases in the sequence, allows various metrics to be easily calculated, and exposes functions for formatting, transforming, and manipulating sequences.

Most operations in the sequence editor operate on a set of strands. Each separate line is assumed to be on a separate strand. Most operations will replace the selection with the indicated transformation (e.g. selecting “Transform”, then “Reverse” will replace the selected sequence with its bases in reverse order). For operations in the “Transform” menu, this can be disabled by unchecking “Replace Selection”; this will cause the transformation to be inserted below each strand.

To view details about each of the menu items in the sequence editor, just hover your mouse over it; help will appear in a tooltip.

9.2 Sequence Designers

- [Web DD](#) - A stochastic, domain-based sequence designer; can be used to design sequences for large systems very quickly by designing a set of noninteracting domains and then threading those domains together to form full strands.
- NUPACK Thermodynamic sequence designer - Uses Caltech’s NUPACK web server to perform multi-objective thermodynamic sequence design.

NUPACK produces thermodynamically optimized sequences, but can be slow for some systems. To use: Enter a design using the NUPACK multi-objective sequence design script, select relevant parameters, and click “Design”; the task will be submitted to the Caltech server, and a popup window will be opened taking you to the results page.

- Multisubjective sequence designer (coming soon)

9.3 Analysis and Simulation

NOTE: The below features are deprecated. Please see [Simulation and Analysis](#) instead.

Workbench provides an interface to the powerful nucleic acid computation package [NUPACK](#). To perform NUPACK calculations of sequences, from the sequence editor, select a set of strands, then select one of the available calculations from the “Compute” menu:

- MFE Complexes - Computes the structure, free energy, and relative concentrations of the various complexes expected at equilibrium. To speed up computation, you can select a maximum complex size to prevent NUPACK from comparing complexes containing more than the indicated number of species.
- Pairwise MFE Complexes - Same as MFE, except the max complex size defaults to 2.
- Subsets MFE - not implemented
- Brute force - not implemented

Chapter 10

Sequence Editor

TBD.

See [Working with sequences](#) for now.

Chapter 11

Thermodynamic simulations with NUPACK

11.1 MFE Complexes

11.2 Pairwise Complexes

Chapter 12

Simulation and Analysis

12.1 NUPACK partition function calculation

Uses Caltech's NUPACK web server to compute the minimum free energy secondary structure of a strand or set of strands. See help from Caltech [here](#).

12.2 Mfold partition function calculation and MFE structure determination

Uses the University of Albany's DINAMelt web server to compute the minimum free energy secondary structure and base pair probabilities. See help from University of Albany [here](#).

12.3 Vienna RNAfold partition function calculation and MFE structure determination

Uses TBI Vienna's RNAfold Websuite to compute the minimum free energy secondary structure of single DNA or RNA strands. See help from TBI [here](#).

Chapter 13

Domain-level reaction enumeration

Reaction enumeration is the process of calculating all possible reactions between, and intermediate complexes formed, from an initial set of complexes. The process proceeds essentially as follows: we begin with a pool of starting complexes, then calculate all possible “fast” (unimolecular) reactions within this pool. Any complexes (or cycles of complexes) without any outgoing “fast” reactions is deemed a “resting complex;” all other species are designated “transient.” Subsequently, “slow” (bimolecular or higher-arity) reactions are then computed between these resting complexes. In DyNAMiC, this process is performed between *domains* (known elsewhere within DyNAMiC as *segments*)—continuous regions of complementarity.¹ That is, rather than considering each distinct base pairing as a different secondary structure, reactions are only considered between longer regions of nucleotides. Pseudoknotted intermediates are not considered.

Once reactions have been enumerated, the enumerator can also generate a “condensed” network of reactions; this network groups all “resting states” (individual resting complexes or cycles of resting complexes connected by fast reactions), and shows only reactions between these resting states. Such a network helps visualize more complex reaction networks by occluding some detail.

13.1 Preparing input for the enumerator

The enumerator accepts several text-based input formats:

¹In the [Nodal](#) compiler and formalism, a distinction is drawn between *segments* (continuous regions of complementarity), and *domains* (groups of segments with a particular behavioral function, such as input or output). In this sense, the reaction enumerator operates on *segments*; here we use the terms *domain* and *segment* interchangeably to refer to a continuous region of complementary bases.

- *Standard input format* (`.enum`) – This is a simple format that is specific to the enumerator. A simple example of the format is included below. The format has three types of statements:
 - **domain** statements declare individual domains, as follows: **domain name : specification**, where:
 - * **name** is the name of the domain (e.g. **a**, **1**, **th**, etc.)
 - * **specification** is either the length of the domain (e.g. a number of bases, or just **long** or **short**) or a sequence (e.g. **NNNNNNN** or **ATTACG** or even a mixture of specific and degenerate bases **AANATCY**)
 - **strand** statements group domains into strands, as follows: **strand name : domains**, where:
 - * **name** is the name of the strand
 - * **domains** is a space-separated list of domains
 - **complex** statements group strands into complexes and assign them a secondary structure, as follows:


```
complex name :
strands
structure
```

 where:
 - * **name** is the name of the complex
 - * **strands** is a space-separated list of strands
 - * **structure** is a domain-wise description of the structure in dot-parenthesis notation
- [Pepper Intermediate Language](#) (`.pil`) – PIL is a general-purpose format for describing domain-level secondary structures. See the [PIL](#) page for details.

Input files for the enumerator can also be generated automatically by the [graphical DIL editor](#).

Here is a simple example of the standard input format:

```
# This file describes the catalytically generated 3 arm junction
# described in Yin et. al. 2008

domain a : 6
domain b : 6
domain c : 6
domain x : 6
domain y : 6
```

```

domain z : 6

strand I : y* b* x* a*
strand A : a x b y z* c* y* b* x*
strand B : b y c z x* a* z* c* y*
strand C : c z a x y* b* x* a* z*

complex I :
I
....

complex A :
A
.(((...)))

complex B :
B
.(((...)))

complex C :
C
.(((...)))

complex ABC :
A B C
(((((((. + )))(((((. + )))))))).

```

13.2 Running the enumerator

Once [an input file has been created](#), you can use the “Run enumerator” button to invoke the enumerator. Click the arrow next to the button to select a specific input format or set advanced options:

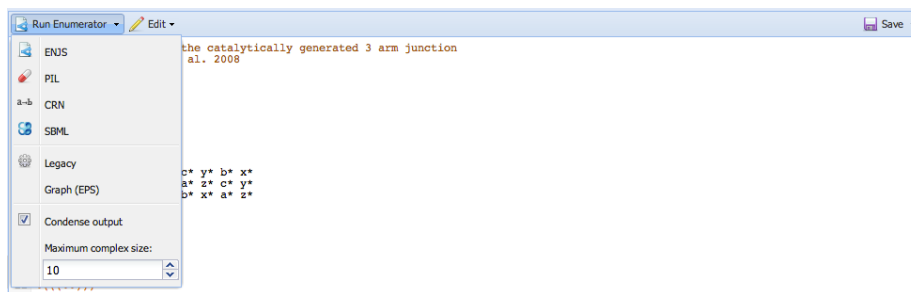


Figure 13.1: Run enumerator button and menu

There are 6 output formats available; you may click any of these menu items to generate output in that format, which will open in a new window:

- Graphical results (ENJS) – produces a file which can be rendered into a graphical, interactive network by DyNAMiC and exported to SVG. This file is also a valid [JSON](#) file and may be suitable for consumption by other tools.
- Pepper Intermediate Language (PIL) – produces a representation of the network, including reactions, in the [Pepper Intermediate Language](#)
- Chemical Reaction Network (CRN) – produces a list of simple reactions between chemical species
- Systems Biology Markup Language (SBML) – produces a representation using the Systems Biology Markup Language, an industry standard format for modeling biological and chemical networks. SBML can be consumed by a reaction simulator, such as [COPASI](#)
- Legacy – produces output in the format of Brian Wolfe’s old enumerator
- Graph (EPS) – produces an EPS file showing the reaction network, laid out using [Graphviz](#)

In addition, the following options may be set:

- Condense output – Check this box to generate the condensed version of the reaction graph. For the ENJS output format, both the condensed network and the full network of reactions will be viewable; for all other output formats, only the condensed network will be written.
- Maximum complex size – Select a maximum number of complexes (beyond which the enumeration should be truncated); this allows systems e.g. with potentially infinite polymers to be enumerate in finite time.

13.3 Viewing enumeration results

Once enumeration is completed, results should open in a new tab automatically. For the text-based output formats (PIL, CRN, SBML, etc.), this tab will show a text editor. For the graphical output format (ENJS), you should see an interactive, graphical representation of the network:

- Complexes are represented by rectangles, while reactions are represented by colored circles (*nodes*). Complexes and reactions are connected by grey lines (*links*). Complexes are colored and outlined colored to indicate whether they are initial, resting or transient.
- You can grab and drag the white background to navigate, and use the mouse wheel to zoom in and out.

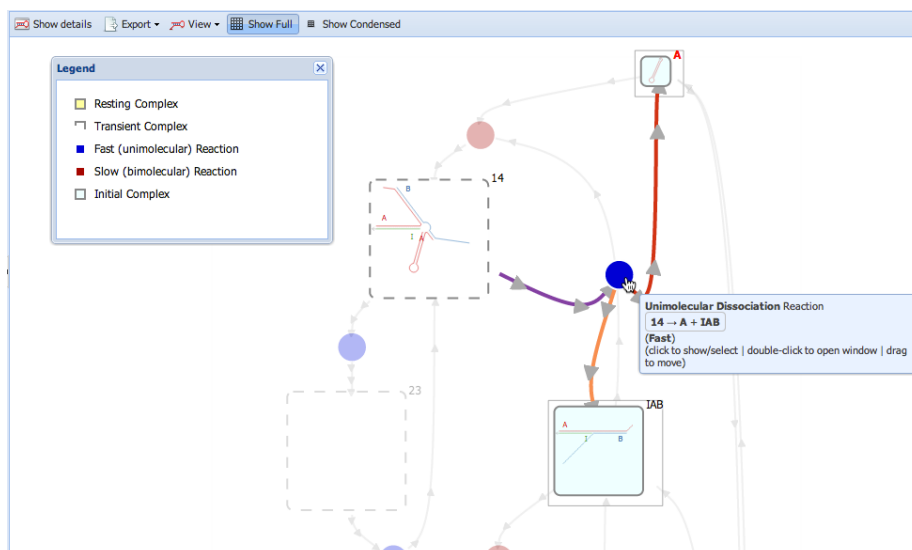


Figure 13.2: Graphical enumeration results

- Mouse-over a complex or reaction to view a description in a tooltip, and to highlight incoming and outgoing links; cool colors show incoming links, while warm colors show outgoing links.
- Click a complex to show its secondary structure; double-click a complex name or a reaction node to view details of the complex or reaction (see [below](#))
- Complexes, reactions, and lines may be dragged to re-arrange the graph for readability.
- Resting states (groups of resting complexes) are outlined in grey

13.3.1 Enumeration results toolbar




Figure 13.3: Enumeration results toolbar

- *Show Details* – if you click to select a complex or reaction, it will be outlined in red; you can then click “Show Details” to view the complex or reaction details (see [below](#))
- *Export* – you can export the entire reaction graph, in its current view, as a [Scalable Vector Graphics](#) (SVG) file which can be viewed or edited in [Adobe Illustrator](#), [Inkscape](#), or similar vector graphics programs.

- *View* – you can choose to view all secondary structures in the network according to various coloring schemes (e.g. by domain, by strand, or by base identity) and view options; for instance, you can choose to view or hide labels for all domains and strands, number bases on the backbone, and view bases as lines, text, or circles (“bubbles”).
- *Show Full* – you can toggle whether the full network of uncondensed reactions is displayed
- *Show Condensed* – you can toggle whether the condensed network of reactions between resting states is displayed

13.3.2 Viewing complex and reaction details

If you double-click on a complex name or a reaction node, a small window will open showing details of the object; this is useful if you would like to compare objects on different parts of the network.




Figure 13.4: Complex details

In complex detail view, you can view the secondary structure of the complex and the constituent strands. You can copy the structure in dot-parenthesis or DU+ notation, or download it as an SVG file. You can apply several view options to the complex as well.

In reaction detail view, you can see a complex preview for each complex in a reaction.




Figure 13.5: Reaction details

13.3.3 Viewing full and condensed reactions

Using the *Show Full* and *Show Condensed* buttons in the toolbar, you can choose to view the full reaction graph, the condensed reaction graph, or both overlaid. For the condensed reaction graph, reaction nodes are larger and shown in a lighter color, while links are darker.




Figure 13.6: Full reaction graph




Figure 13.7: Condensed reaction graph




Figure 13.8: Both reaction graphs overlaid

Chapter 14

Thermodynamic simulations with NUPACK

14.1 MFE Complexes

14.2 Pairwise Complexes

Chapter 15

Customization

Workbench includes a documented [Javascript](#) API which you can use to extend the built-in capabilities, perform custom processing, or modify the default behavior of applications. There are two main ways to do this:

- Console - You can execute arbitrary Javascript within the context of the active Workbench [Application](#) using the console.
- Custom scripts - You can write and execute complete Javascript programs from within Workbench. This is most useful for complicated workflows which need to be automated.

15.1 Console

You can also run small snippets of code within a particular application, using the Workbench console.

To open the console:

- From the main menu, select ‘Tools’, then ‘Console’.
- Enter your code in the console window which appears at the lower right.

15.2 Custom scripts

You can use a custom script to write and execute a Javascript program within Workbench; this allows you access to all of the programmatic tools that Workbench provides, while letting you automate complicated or repetitive behavior.

To build a custom script:

- [Create](#) a new Javascript file.
- Open the file and enter your script.
- Click 'run' to execute your script.

Chapter 16

DyNAMiC Workbench Server

16.1 Overview

Workbench server is the part of the Workbench suite which is responsible for managing computationally intensive tasks, and storing files for Workbench users. There are two ways you might access a Workbench server:

- Hosted installation - Workbench server is intended to be installed on a cluster computing platform and made available via the web. In this case, you need only interact with the web-based Workbench client interface. You'll just need an invitation code from whoever runs the server, and you'll be able to create an account and begin using Workbench. The only hosted installation currently available is at provided by the [Molecular Systems lab at Harvard](#).
- Local installation (hosting your own server) - For testing purposes, you may wish to host your own Workbench server. This requires a bit more effort, but you have full control over the entire system. See below for details about how to set up your own server.

16.2 Server Tools

Workbench ships with several server tools installed. For details, see [Server tools](#).

16.3 Hosting your own server

If you are reading this documentation, you’ve likely already obtained a copy of Workbench from the [Molecular Systems Lab](#). Because of the number of external dependencies that the Workbench server has, and the relative difficulty in setting them up, Workbench server is deployed as a [VirtualBox](#) appliance. VirtualBox is a free virtualization platform provided by Oracle. This means Workbench will run as a Virtual Machine, with its own isolated operating system, file system, process management, etc. Therefore, you only need to install the virtual machine, and you have access to all of the relevant [server tools](#) (such as NUPACK, SpuriousDesign, the Nodal and Pepper compilers, etc.) without needing to configure them individually. This setup has the added benefit that if Workbench or one of its server tools crashes, it won’t affect your host machine.

It’s important to understand how this setup works: The Workbench server virtual machine will run (using VirtualBox) on your computer (which is called the “host” in this circumstance); it contains a separate operating system (the “guest” operating system, which in this case is a version of Ubuntu linux), and a lot of software, including a web server and the server tools. All of this software which will run within the virtual machine, sharing your processor and memory, but essentially isolated from your computer. There are two special communication channels between the virtual machine and the host:

- Shared folders: this VirtualBox feature allows folders on the host to be mirrored in the guest, and vice-versa. This lets you to access your files stored on Workbench from within your normal operating system file manager (e.g. Finder, Nautilus, Windows Explorer).
- Host-to-guest network: this creates a special network only between the host and guest. This means that the virtual machine will not be visible to the internet at large, but it will be able to connect only to the host (for instance to expose the Workbench web server).

16.3.1 Running the server

As part of the [installation](#) process, you’ll install a copy of the pre-packaged Workbench Server virtual machine on your computer. You’ll be able to launch the virtual machine (VM) directly from the VirtualBox desktop application, or using the Workbench Server manager on Mac OS X.

The actual server component is configured to launch automatically when the appliance starts. That means if you just use the Workbench interface, all you need to do is launch the appliance and point your browser

However, if you want to tweak the server beyond what’s described in the [customization](#) page, or to use any of the installed server tools directly (from the command line), you’ll need to log in.

16.3.2 Logging in to the server

To log in to the server, you use a separate set of user credentials (different from the username and password that you use to log into your host machine, or that you use to log in to the Workbench client interface on the web). These credentials are preset when you download Workbench, although you're encouraged to change them.

The predefined credentials are:

```
Username: 'webserver-user'  
Password: ' '
```

(single quotes are not part of the username or password; the password is a single space: ' ').

The recommended method for logging in to the virtual machine is via [SSH](#). This will allow you command-line access to the server.

To connect to the server via SSH:

- On Mac OS X or Linux, open a Terminal, and enter the following command: `ssh webserver-user@192.168.56.10`. You will be prompted to enter `webserver-user@192.168.56.10's password:`; enter the password.
- On Windows: you'll need to download an SSH client, such as [PuTTY](#). Open your SSH client, and login using credentials like this:

```
host: 192.168.56.10  
port: 22  
user: (see above)  
password: (see above)
```

16.3.3 Using the web interface

The server will start automatically after the VM has finished booting (your server should be running by the time you see a login prompt). To view the web interface, point your web browser to: `[http://192.168.56.10:3000/]`.

See [documentation](#) for the web interface.

16.3.4 Interacting with the server via SSH

Once you’ve logged in to the server with SSH, if you’re comfortable, you can play around with shell access to the server.

The actual server process is described in a shell script: `~/startup`, which you can look at if you’re curious. This script starts the [Node JS web server](#) and the [Mongo database](#) processes, which do the heavy lifting of running the server. When `startup` is killed, it intelligently kills both processes.

`startup` is in turn controlled by an [Upstart](#) script, located in `/etc/init/workbench.conf`. The upstart script makes sure that the server gets launched on startup, killed on shutdown, and restarted if it crashes. You can control the server using Upstart commands:

- `sudo start workbench` – starts the server
- `sudo stop workbench` – stops the server
- `sudo status workbench` – tells you if the server is running or not

One other shell script is provided for your convenience: `~/repair`. Occasionally, the database server doesn’t shut down properly (this happens when the virtual machine is powered off without killing the server process). If when you launch the server normally and attempt to log in via the web interface, the login progress bar just keeps resetting, your database needs to be repaired; in that case, run:

```
sudo stop workbench
sh ~/repair
sudo start workbench
```

Note: `sudo` is required because administering Upstart processes requires administrator privileges. However, `~/startup` is actually run as `webserver-user`. `webserver-user` is currently on the `sudoers` list, but the plan is to eventually create a separate user account for administration and return `webserver-user` to limited privileges again.

16.3.5 Shutting down the server

To shut down the server and avoid damaging the database, simply shut down the virtual machine by:

- Closing its application window (titled “VirtualBox VM” or some such thing), and selecting “Send the shutdown signal”, or:
- Entering: `VBoxManage controlvm 'DyNAMiC Workbench Server (0.3.0)' acpipowerbutton` on the command line in your host operating system (not the VM or ssh), or
- Entering `sudo shutdown 0` on the command line in the VM

16.3.6 Server Manager Application

On Mac OS X, a more intuitive application has been provided to automate some of these tasks; you can launch it by opening the ‘Workbench Server’ application in your Applications folder. You’ll still need to use SSH to fix something if it breaks (to repair the database, for instance). However, you can start and stop the virtual machine safely by using the “Start” and “Stop” buttons, and you can do other convenient things like open the client interface in a web browser.

To start the server from the Server Manager application:

- Click ‘Start’ from the server control window. You will see a VirtualBox window open, displaying the screen for your Virtual Machine. Wait until you see a command prompting you to log in.
- You can choose to log in and interact with the server (as described above), or you can just go directly to the [web interface](#).

Chapter 17

Scripting

17.1 Overview

17.2 Console

17.3 Scripting Applications

Chapter 18

Workbench Application Development

Coming soon.

18.1 Overview

18.2 Bundling

18.3 API Documentation